

Module 4.0

Imagine that we are zipping two tensors `a` and `b` of shape $(I, 1)$ and $(1, J)$ where I is very large and J is a small fixed size. (this is a special case of the zip in the homework)

```
def specialized_zip_cuda(out, a, b):  
  
    x = numba.cuda.blockIdx.x * THREADS + numba.cuda.threadIdx.x  
    thread_x = numba.cuda.threadIdx.x  
  
    shared = numba.cuda.shared.array(J)  
    if x < I:  
  
        if thread_x < J:  
  
            shared[thread_x] = b[thread_x]  
  
        numba.cuda.syncthreads()  
  
        local = a[x]  
  
        for j in range(J):  
  
            out[x, j] = local * shared[j]
```

1

1 point



Imagine that $I = 1000$ and $J = 100$, what is the maximum number of global *reads* there from any thread?

2

1 point



Imagine that $I = 1000$ and $J = 100$, how many shared *reads* are there from each thread?

3

1 point



Imagine that $I = 1000$ and $J = 100$, what is the max number of shared memory *writes* from any thread?